



## Phase 1: Offline Uncertainty Estimation No training cost



## Phase 2: Online Uncertainty-Calibrated Training On-policy loop



Infrastructure: 4x H100, Ray + DeepSpeed + vLLM