
Uncertainty-Calibrated On-Policy Distillation for Large Language Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

On-policy distillation trains a student on its own outputs with teacher feedback to avoid the train-inference distribution mismatch present in standard distillation. However, a fundamental underexplored challenge is that the teacher itself is imperfect; consequently, naively imitating erroneous teacher outputs propagates mistakes into the student. Prior work such as Self-Distilled Reasoner [1] mitigates this by injecting ground-truth chain-of-thought solutions into the teacher’s system prompt to guarantee the teacher with correct reasoning. But this approach depends on labeled data that is unavailable in many practical settings. We introduce **UOPD** (Uncertainty-calibrated On-Policy Distillation), a framework that addresses teacher errors without ground-truth supervision. Teacher reliability is estimated offline by computing its **uncertainty**, and the distillation objective is reweighted inversely proportional to the teacher’s uncertainty with no additional training cost. We further release **SURE-Math**, a dataset annotated with teacher semantic entropy scores. We evaluate UOPD across mathematical reasoning, summarization, translation, and instruction tuning tasks. UOPD converges faster and achieves state-of-the-art results compared to both GRPO and standard on-policy distillation.

17 **1 Introduction**

18 **2 Preliminary**

19 **2.1 On-Policy Knowledge Distillation**

20 Standard knowledge distillation [2] trains a student model $p_S(\cdot; \theta_S)$ to minimize the KL divergence
21 from a frozen teacher p_T on a fixed dataset \mathcal{D} :

$$\mathcal{L}_{\text{KD}}(\theta_S) = \mathbb{E}_{x \sim \mathcal{D}} [\text{KL}(p_T(\cdot | x) \| p_S(\cdot | x; \theta_S))]. \quad (1)$$

22 This offline objective suffers from a train-inference distribution mismatch: during training the
23 student observes token-level distributions conditioned on dataset prefixes, whereas at inference it
24 must generate from its own previously predicted tokens. On-policy distillation [3] resolves this by
25 generating training sequences from the student itself:

$$\mathcal{L}_{\text{OPD}}(\theta_S) = \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y \sim p_S(\cdot | x; \theta_S)} \left[\sum_{t=1}^{|y|} \text{KL}(p_T(\cdot | x, y_{<t}) \| p_S(\cdot | x, y_{<t}; \theta_S)) \right]. \quad (2)$$

26 Since the training distribution now matches the inference distribution, on-policy distillation substantially
27 reduces compounding errors in chain-of-thought reasoning [4].

28 **2.2 Uncertainty in LLMs**

29 Uncertainty in language model outputs can be decomposed into two fundamentally distinct sources
30 [5]. **Aleatoric uncertainty** is irreducible and arises from the inherent ambiguity of natural language,
31 where the same meaning can be expressed in infinitely many surface forms. Consider the prompt
32 “*What is the capital of the UK?*” Given this question, the responses “*London*”, “*London is the capital*
33 *of the UK*”, and “*The UK’s capital is London*” are semantically identical, while their token-level
34 divergence is large. A teacher that consistently produces such paraphrases is not making errors; it
35 is exhibiting natural lexical variation. Penalising this variation at the token level conflates surface
36 diversity with genuine unreliability.

37 **Epistemic uncertainty**, by contrast, reflects the teacher’s *lack of knowledge*. It arises when the
38 teacher generates contradictory answers across samples, indicating that it does not reliably know the
39 correct response to a given prompt. Standard token-level objectives cannot distinguish these two
40 sources, because both manifest as high distributional variance in the output space.

41 Semantic entropy [6] resolves this ambiguity by operating over *meaning* rather than surface tokens.
42 Responses are first clustered into semantic equivalence classes \mathcal{C} , and entropy is computed over the
43 resulting distribution

$$H_{\text{sem}}(x) = - \sum_{c \in \mathcal{C}} p(c | x) \log p(c | x), \quad p(c | x) \propto \sum_{y \in c} p_T(y | x). \quad (3)$$

44 Paraphrases of the same answer collapse into a single class, suppressing aleatoric noise. $H_{\text{sem}}(x)$
45 is therefore a faithful, label-free measure of epistemic uncertainty; it is high only when the teacher
46 genuinely disagrees with itself across semantically distinct answers.

47 **3 Method**

48 We propose UOPD, a two-phase framework for uncertainty-calibrated on-policy distillation. In the
49 first phase, we estimate the teacher’s epistemic uncertainty for each training prompt by computing
50 semantic entropy over multiple teacher rollouts. This computation is performed entirely offline before
51 training begins, producing a per-prompt uncertainty weight that is stored alongside the training data.
52 In the second phase, the student is trained on-policy with token-level distillation from the teacher,
53 where each sample’s contribution to the loss is calibrated by its precomputed uncertainty weight.

54 The key advantage of this design is the separation of uncertainty estimation from training. Because
55 semantic entropy is computed once before the training loop, UOPD introduces zero additional
56 overhead during training while still enabling uncertainty-aware distillation. Figure 1 illustrates the
57 overall pipeline. We describe each component in detail below.

58 **3.1 Offline Semantic Entropy Estimation**

59 Given a training set of prompts $\{x_i\}_{i=1}^M$, we first estimate the teacher’s epistemic uncertainty per
60 prompt. For each prompt x_i , we sample N responses from the teacher $\{y_1^T, \dots, y_N^T\} \sim p_T(\cdot | x_i)$,
61 recording each response’s sequence log-probability $\log p_T(y_n^T | x_i)$ and token count $|y_n^T|$.

62 **Semantic clustering.** As discussed in Section 2, responses such as “*London*”, “*London is the*
63 *capital of the UK*”, and “*The UK’s capital is London*” are semantically identical but produce large
64 token-level divergence. To prevent this aleatoric variation from inflating uncertainty estimates, we
65 extract the final answer from each teacher response and cluster semantically equivalent answers into
66 equivalence classes $\mathcal{C} = \{c_1, \dots, c_J\}$. Two answers are placed in the same cluster only if one can be
67 derived from the other, i.e., they express the same meaning in different surface forms (e.g., $\frac{1}{2}$ and 0.5,
68 or $x = 3$ and 3). Answers that match as exact strings are merged directly. For non-matching pairs, a
69 Qwen-4B judge [7] determines pairwise semantic equivalence, and all judgments are resolved into
70 clusters with a Union-Find algorithm.

71 **Intra-cluster vs. inter-cluster entropy.** This clustering decomposes the total entropy of teacher
72 outputs into two interpretable components. *Intra-cluster entropy* captures the lexical diversity *within*
73 a semantic equivalence class. A cluster containing “0.5”, “ $\frac{1}{2}$ ”, and “the answer is one half” has high
74 intra-cluster entropy, reflecting the teacher’s expressive richness rather than genuine confusion. This
75 variation is aleatoric and should not be penalized. *Inter-cluster entropy*, by contrast, measures the

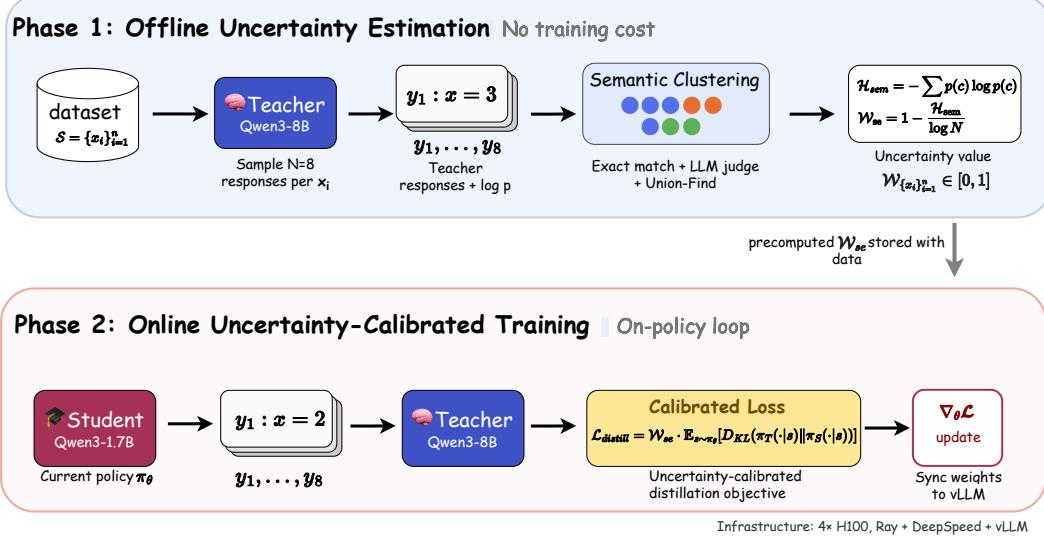


Figure 1: Overview of UOPD. Phase 1 estimates teacher reliability offline by computing semantic entropy over meaning clusters rather than surface-form token distributions. Phase 2 performs on-policy distillation whose per-sample loss is calibrated by the precomputed uncertainty weights. Because the two phases are fully decoupled, uncertainty estimation introduces zero additional overhead during training.

76 spread of probability mass *across* semantically distinct answer classes. When the teacher assigns
 77 substantial mass to multiple contradictory clusters (e.g., both “3” and “5” for the same prompt),
 78 inter-cluster entropy is high, signaling epistemic uncertainty. Our semantic entropy H_{sem} captures
 79 precisely the inter-cluster component by collapsing all within-cluster variation before computing
 80 entropy.

81 **Probability-weighted semantic entropy.** Each response’s log-probability is first length-normalized
 82 to avoid penalizing longer outputs. The probability mass of semantic class c is then computed as

$$p(c | x_i) = \frac{\sum_{n \in c} \exp(\log p_T(y_n^T | x_i) / |y_n^T|)}{\sum_{n=1}^N \exp(\log p_T(y_n^T | x_i) / |y_n^T|)}, \quad (4)$$

83 and the semantic entropy is

$$H_{\text{sem}}(x_i) = - \sum_{c \in C} p(c | x_i) \log p(c | x_i). \quad (5)$$

84 We normalize by the maximum possible entropy to obtain a value in $[0, 1]$ and define the per-prompt
 85 distillation weight as

$$w_{\text{se}}(x_i) = 1 - \frac{H_{\text{sem}}(x_i)}{\log N}. \quad (6)$$

86 A weight near 1 indicates that the teacher consistently agrees on the same semantic answer (low
 87 epistemic uncertainty), while a weight near 0 indicates contradictory responses across clusters (high
 88 epistemic uncertainty). These weights are precomputed and stored with the training data.

89 3.2 Uncertainty-Calibrated On-Policy Distillation

90 At each training step, the student generates K responses per prompt from its current policy
 91 $\{y^{(1)}, \dots, y^{(K)}\} \sim p_S(\cdot | x; \theta_S)$. For each student-generated response y , the frozen teacher pro-
 92 vides its top- k logit values and corresponding token indices at every position, reducing memory and
 93 communication from the full vocabulary size V to $k \ll V$.

94 The token-level distillation loss at position t is the cross-entropy between the teacher’s and student’s
 95 distributions over the top- k tokens

$$\ell_t = - \sum_{j=1}^k \tilde{p}_T(v_j | x, y_{<t}) \log p_S(v_j | x, y_{<t}; \theta_S), \quad (7)$$

96 where \tilde{p}_T denotes the teacher’s probability renormalized over the top- k tokens and $\{v_1, \dots, v_k\}$ are
 97 the teacher’s top- k token indices.

98 The standard on-policy distillation objective averages ℓ_t uniformly over all response tokens, treating
 99 every prompt equally regardless of teacher reliability. UOPD instead calibrates each sample’s
 100 contribution by its precomputed semantic entropy weight $w_{se}(x)$. For a prompt x and student-
 101 generated response y , the uncertainty-calibrated loss is

$$\mathcal{L}(x, y; \theta_S) = \frac{w_{se}(x)}{\sum_t \mathbb{1}[t \in \text{resp}]} \sum_{t \in \text{resp}} \ell_t, \quad (8)$$

102 where the sum runs over response token positions (excluding the prompt). When the teacher is
 103 confident about a prompt ($w_{se} \approx 1$), the student receives the full distillation signal. When the teacher
 104 is uncertain ($w_{se} \approx 0$), the loss is suppressed, preventing the propagation of erroneous teacher
 105 guidance.

106 3.3 Overall Training Objective

107 The full training objective averages over all prompts and their K on-policy rollouts

$$\mathcal{L}_{\text{UOPD}}(\theta_S) = \frac{1}{|\mathcal{D}|} \sum_{x \in \mathcal{D}} \frac{1}{K} \sum_{j=1}^K \mathcal{L}(x, y^{(j)}; \theta_S). \quad (9)$$

108 At each training step, the updated student weights are synchronized back to the vLLM generation
 109 engines, ensuring that the next round of rollouts reflects the latest policy. This on-policy loop
 110 continues until convergence.

111 Algorithm 1 summarizes the complete UOPD pipeline. Phase 1 (lines 2–7) iterates over all training
 112 prompts, sampling N teacher responses per prompt, clustering them into semantic equivalence classes,
 113 and storing the resulting uncertainty weight w_{se} . Phase 2 (lines 9–14) performs the on-policy training
 114 loop: at each step, the student generates K rollouts, queries the teacher for top- k logits at every token
 115 position, and updates its parameters with the uncertainty-weighted distillation loss, after which the
 116 refreshed weights are synchronized to the generation engines for the next iteration.

117 4 Experiments

118 4.1 SURE-Math Dataset

119 We introduce **SURE-Math** (Semantic-Uncertainty REasoning Math), a mathematical reasoning
 120 dataset in which every problem is annotated with the teacher’s precomputed semantic entropy score.

121 We first collect 1,000 mathematics problems spanning middle-school to competition level, each paired
 122 with a verified ground-truth answer. Of these, 200 are reserved as a held-out test set (Section 4.3); the
 123 remaining 800 are added to the training pool.

124 We further aggregate seed problems from public sources across four difficulty tiers. *Easy*: ScaleQuest-
 125 Math [8] (15K). *Medium*: NuminMath-CoT [9] (10K) and MATH [10] (12.5K). *Hard*: Omni-MATH
 126 [11] (4.4K) and OlympiadBench [12] (5K). *Competition*: AIME 2024–2025, HMMT Feb & Nov
 127 2025, and AMO-Bench [13].

128 Starting from these seeds, we synthesize new problems with Qwen2.5-72B-Instruct using an Evol-
 129 Instruct [14] style pipeline. Six evolution strategies are applied, each realized by a dedicated
 130 system prompt: *harder* (increase reasoning steps or add constraints), *rewrite* (change context and
 131 wording while preserving the underlying skill), *algebraize* (replace concrete values with variables),
 132 *apply* (embed the concept in a real-world scenario), *compose* (combine with a different branch of

Algorithm 1 Uncertainty-Calibrated On-Policy Distillation (UOPD)

Require: Training prompts $\mathcal{D} = \{x_i\}_{i=1}^M$, teacher p_T , student $p_S(\cdot; \theta_S)$, SE samples N , rollouts K , top- k

1: — **Phase 1: Offline Semantic Entropy (Sec. 3.1) —**

2: **for** each prompt $x_i \in \mathcal{D}$ **do**

3: Sample N responses $\{y_1^T, \dots, y_N^T\} \sim p_T(\cdot | x_i)$

4: Extract final answers; cluster into semantic classes \mathcal{C}

5: Compute semantic entropy $H_{\text{sem}}(x_i)$ (Eq. 5)

6: Store weight $w_{\text{se}}(x_i) = 1 - H_{\text{sem}}(x_i)/\log N$ (Eq. 6)

7: **end for**

8: — **Phase 2: On-Policy Distillation (Sec. 3.2) —**

9: **while** not converged **do**

10: Sample mini-batch $\mathcal{B} \subset \mathcal{D}$

11: **for** each $x \in \mathcal{B}$ **do**

12: Generate K rollouts $\{y^{(1)}, \dots, y^{(K)}\} \sim p_S(\cdot | x; \theta_S)$

13: **for** each rollout $y^{(j)}$ **do**

14: Query teacher for top- k logits at each token position of $y^{(j)}$

15: Compute token-level CE: $\ell_t = -\sum_v \tilde{p}_T(v) \log p_S(v; \theta_S)$

16: **end for**

17: Compute weighted loss $\mathcal{L}(x, y^{(j)}; \theta_S)$ (Eq. 8)

18: **end for**

19: Update θ_S with $\nabla_{\theta_S} \frac{1}{|\mathcal{B}|} \sum_{x \in \mathcal{B}} \frac{1}{K} \sum_{j=1}^K \mathcal{L}(x, y^{(j)}; \theta_S)$

20: Sync θ_S to vLLM generation engines

21: **end while**

133 mathematics), and *competition* (transform into AMC/AIME/Olympiad style). Each seed undergoes
134 1–5 rounds of iterative evolution depending on its difficulty tier: easy seeds receive 1 round, medium
135 seeds 2 rounds, hard seeds 3 rounds, and competition seeds 5 rounds. In follow-up rounds, a different
136 strategy (drawn from *harder*, *competition*, and *compose*) is applied to the output of the previous
137 round, progressively increasing difficulty. Strategy selection within each tier is weighted to favor
138 *harder* and *competition* for higher-tier seeds.

139 For each of the 8,000 problems, we sample $N=8$ responses from Qwen3-8B [7] at temperature 0.7,
140 extract the \boxed{} answer from each response, and compute the probability-weighted semantic
141 entropy as described in Section 3.1. The resulting per-problem semantic entropy score and the
142 corresponding distillation weight w_{se} are stored alongside each problem.

143 We reserve 200 of the 1,000 manually collected problems (which carry ground-truth labels) as a
144 held-out test set. The remaining 800 curated problems and all 8,000 evolved problems form the
145 training set.

146 **4.2 Experimental Setup**

147 **Models.** We evaluate three student–teacher pairs of increasing scale: Qwen3-1.7B/Qwen3-8B,
148 Qwen3-4B/Qwen3-8B, and Qwen3-8B/Qwen3-30B [7], all instruction-tuned variants. Semantic
149 clustering for the offline SE computation uses a Qwen2.5-3B-Instruct judge for pairwise semantic
150 equivalence.

151 **Training details.** We train with the on-policy distillation objective described in Eq. 8. At each step,
152 the student generates $K=4$ rollouts per prompt. The teacher provides its top-512 logit values at every
153 token position. We use a learning rate of 3×10^{-6} with a cosine schedule and 5% warmup, AdamW
154 with $(\beta_1, \beta_2) = (0.9, 0.95)$, gradient clipping at 1.0, and a global batch size of 128. Training runs for
155 one epoch over 50,000 samples (with replacement from the 8,800 training problems). All models use
156 bfloat16 precision and FlashAttention-2.

157 For each student–teacher pair, we compare UOPD against the following methods. (1) **Base**: the
158 student model without any distillation (lower bound). (2) **Teacher**: the teacher model (upper
159 bound). (3) **Standard OPD**: on-policy distillation with uniform weighting ($w_{\text{se}} = 1$ for all prompts),

Table 1: Main results on mathematical reasoning benchmarks. We report pass@1 (greedy) and avg@16 accuracy (%). Best student-sized results are **bolded**.

Method	MATH-500		AIME 2025		AMO-Bench		SURE-Math	
	pass@1	avg@16	pass@1	avg@16	pass@1	avg@16	pass@1	avg@16
Student: Qwen3-1.7B Teacher: Qwen3-8B								
Base (1.7B)	–	–	–	–	–	–	–	–
Standard OPD	–	–	–	–	–	–	–	–
GRPO	–	–	–	–	–	–	–	–
Self-Distilled Reasoner	–	–	–	–	–	–	–	–
UOPD (Ours)	–	–	–	–	–	–	–	–
Student: Qwen3-4B Teacher: Qwen3-8B								
Base (4B)	–	–	–	–	–	–	–	–
Standard OPD	–	–	–	–	–	–	–	–
GRPO	–	–	–	–	–	–	–	–
Self-Distilled Reasoner	–	–	–	–	–	–	–	–
UOPD (Ours)	–	–	–	–	–	–	–	–
Student: Qwen3-8B Teacher: Qwen3-30B								
Base (8B)	–	–	–	–	–	–	–	–
Standard OPD	–	–	–	–	–	–	–	–
GRPO	–	–	–	–	–	–	–	–
Self-Distilled Reasoner	–	–	–	–	–	–	–	–
UOPD (Ours)	–	–	–	–	–	–	–	–

160 following [3]. (4) **GRPO**: Group Relative Policy Optimization, which trains the student using reward
 161 signals from correct and incorrect rollouts. (5) **Self-Distilled Reasoner**: on-policy distillation with
 162 ground-truth chain-of-thought injected into the teacher’s system prompt [1].

163 4.3 Evaluation

164 **Benchmarks.** We evaluate all methods on three public mathematical reasoning benchmarks of
 165 increasing difficulty, plus our held-out test set: **MATH-500** [10], 500 competition-level problems
 166 spanning seven subjects; **AIME 2025**, 30 problems from the American Invitational Mathematics
 167 Examination (Parts I and II); **AMO-Bench** [13], 50 IMO-level competition problems; and our
 168 held-out **SURE-Math test set** (200 curated problems with ground-truth labels).

169 **Metrics.** We report pass@1 accuracy with greedy decoding (temperature 0) and avg@16 accuracy
 170 with 16 samples at temperature 1.2 and top- $p=0.95$.

171 4.4 Main Results

172 Table 1 presents the main comparison across all benchmarks.

173 4.5 Ablation Studies

174 **Effect of uncertainty weighting.** We compare three weighting strategies: (1) uniform weighting
 175 ($w_{se} = 1$), which reduces to standard on-policy distillation; (2) binary filtering, which discards all
 176 prompts with $H_{sem} > \tau$ for a threshold τ ; and (3) soft weighting ($w_{se} = 1 - H_{sem}/\log N$), which is
 177 our default. Table 2 reports the results.

Table 2: Ablation on uncertainty weighting strategies.

Weighting	MATH-500	AIME 2025	AMO-Bench	SURE-Math
Uniform ($w = 1$)	–	–	–	–
Binary filter ($\tau = 0.3$)	–	–	–	–
Binary filter ($\tau = 0.5$)	–	–	–	–
Soft weighting (UOPD)	–	–	–	–

178 **Number of SE samples** N . The number of teacher rollouts N used for semantic entropy estimation
179 controls the resolution of the uncertainty estimate. We vary $N \in \{2, 4, 8, 16\}$ and measure
180 downstream distillation performance.

181 **Top- k logit truncation.** Transmitting the full vocabulary ($|V| = 151,936$) is memory-intensive.
182 We compare $k \in \{128, 512, 2048\}$ and full-vocabulary distillation, measuring both accuracy and
183 peak GPU memory.

184 **Number of student rollouts** K . We ablate $K \in \{1, 2, 4, 8\}$ on-policy rollouts per prompt to
185 understand the trade-off between training diversity and computational cost.

186 4.6 Analysis

187 Convergence speed.

188 **Semantic entropy distribution.** Figure ?? shows the distribution of semantic entropy scores across
189 SURE-Math. The majority of problems have low SE, indicating that the teacher is confident on most
190 prompts. The long tail of high-SE problems represents cases where the teacher genuinely disagrees
191 with itself, and these are precisely the samples that UOPD downweights.

192 **Validating SE as a teacher error signal.** A core assumption of UOPD is that high semantic
193 entropy indicates problems the teacher is likely to answer incorrectly. We verify this directly on the
194 SURE-Math training set. For each of the 16,330 training problems, we obtain ground-truth labels
195 by generating solutions with Qwen2.5-72B-Instruct and then evaluate the teacher (Qwen3-8B) with
196 greedy decoding on the same problems. The teacher’s predicted answer is compared against the
197 72B-verified ground truth using exact-match and symbolic equivalence checking.

198 Figure ??(a) bins the training problems by their precomputed SE score and plots the teacher’s
199 accuracy within each bin. The relationship is strikingly monotonic: for problems with near-zero SE
200 (< 0.1), the teacher achieves 87% accuracy, whereas for problems with $SE > 0.9$, accuracy drops
201 to just 12%. Figure ??(b) treats teacher error as a binary classification target and uses SE as the
202 predictor, yielding an AUROC of **0.7949**. This confirms that probability-weighted semantic entropy
203 is a reliable proxy for teacher correctness, validating the use of $w_{se} = 1 - \bar{H}_{sem}$ as a distillation
204 weight: samples where the teacher is most likely wrong receive the lowest weight, while confident
205 and correct samples dominate the training signal.

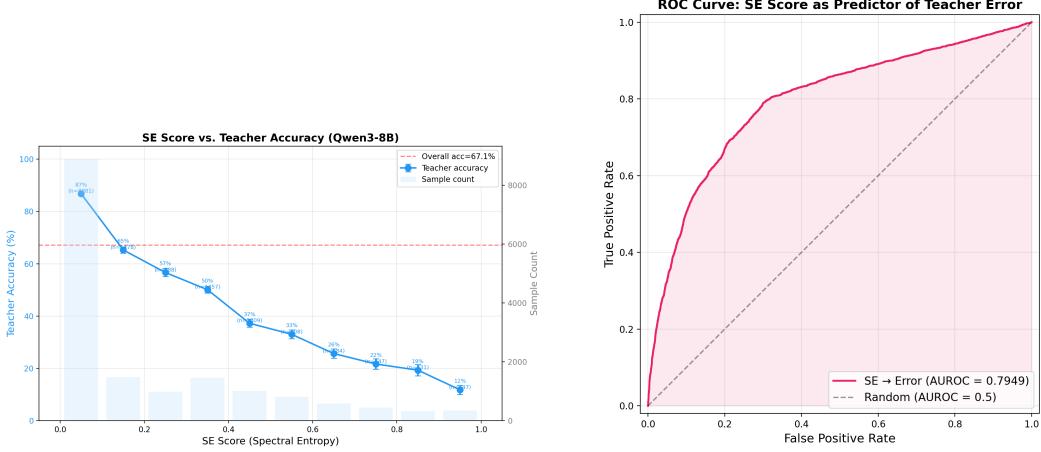
206 **Qualitative examples.** Table ?? shows representative examples where UOPD’s uncertainty weighting
207 helps. For low-SE prompts, the teacher provides consistent guidance and the student learns
208 effectively. For high-SE prompts, the teacher produces contradictory answers; UOPD suppresses
209 these samples, preventing the student from learning incorrect reasoning patterns.

210 5 Related Work

211 5.1 Knowledge Distillation for Language Models

212 Traditional knowledge distillation [2] trains student models to match teacher output distributions.
213 For autoregressive language models, supervised KD [15] and sequence-level KD [16] are widely
214 used. However, these off-policy approaches suffer from distribution mismatch between training
215 (teacher-generated or ground-truth sequences) and inference (student-generated sequences).

216 **On-policy distillation** [3] addresses this by training students on their own generated outputs, using
217 teacher logits as labels. GKD (Generalized Knowledge Distillation) demonstrates strong improve-
218 ments on summarization and translation tasks. Our work extends on-policy distillation to reasoning
219 tasks by incorporating verification signals and contrastive objectives.



(a) Teacher accuracy (Qwen3-8B, greedy) vs. SE score. Accuracy decreases monotonically from 87% to 12% as SE increases.

(b) ROC curve using SE as a predictor of teacher error, achieving an AUROC of 0.7949.

Figure 2: Empirical validation that semantic entropy predicts teacher error. Ground-truth labels are obtained from Qwen2.5-72B-Instruct on the 16,330 SURE-Math training problems. High SE reliably identifies samples where the teacher answers incorrectly, justifying the uncertainty-based weighting in UOPD.

220 5.2 Learning from Verification Feedback

221 Recent work on mathematical reasoning leverages verification to improve model training. GRPO [?] 222 and similar RL-based approaches optimize for verified correctness using policy gradient methods. 223 V-STaR [?] iteratively generates verified solutions for self-improvement.

224 However, these methods rely solely on sparse binary rewards (correct/incorrect) and do not leverage 225 dense token-level teacher guidance. TCD bridges this gap by combining verification with token-level 226 distillation.

227 5.3 Contrastive Learning for Language Models

228 DPO [17] introduced preference-based contrastive learning for alignment, training models to prefer 229 chosen responses over rejected ones. SimPO [?] and other variants explore different contrastive 230 formulations.

231 Our work adapts contrastive learning to the distillation setting: we use verified correct responses 232 (from reference data or teacher generation) as "chosen" and student-generated incorrect responses as 233 "rejected", enabling the student to learn from its mistakes with teacher guidance.

234 6 Conclusion

235 We presented Token-level Contrastive Distillation (TCD), a framework that combines on-policy 236 generation, verification-based learning, and dense teacher guidance for distilling reasoning capabilities 237 into smaller language models. By unifying distillation on correct traces with contrastive learning 238 on errors, TCD effectively leverages both sparse verification signals and rich token-level teacher 239 feedback.

240 Our experiments on GSM8K demonstrate that TCD enables a 1.7B student model to achieve strong 241 mathematical reasoning performance when distilled from an 8B teacher. The framework is efficient 242 and scalable, using vLLM for on-policy generation and memory-optimized techniques for handling 243 large vocabularies.

244 **Limitations and Future Work.** **Scalability:** Current experiments use 4 GPUs; scaling to multi-
245 node setups requires careful optimization of communication backends (NCCL vs. Gloo). **Generaliza-**
246 **tion:** We focus on mathematical reasoning; extending to other reasoning tasks (code, commonsense)
247 is an important direction. **Teacher quality:** Our approach assumes access to a capable teacher;
248 exploring self-improvement scenarios is promising.

249 TCD opens avenues for efficient reasoning model deployment by making strong reasoning capabilities
250 accessible in compact models suitable for resource-constrained environments.

251 References

- 252 [1] Siyan Zhao, Zhihui Xie, Mengchen Liu, Jing Huang, Guan Pang, Feiyu Chen, and Aditya
253 Grover. Self-distilled reasoner: On-policy self-distillation for large language models. *arXiv*
254 preprint arXiv:2601.18734, 2026.
- 255 [2] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network.
256 *arXiv preprint arXiv:1503.02531*, 2015.
- 257 [3] Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu
258 Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-
259 generated mistakes. In *International Conference on Learning Representations (ICLR)*, 2024.
- 260 [4] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, and
261 Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. *Advances*
262 in *Neural Information Processing Systems*, 35:24824–24837, 2022.
- 263 [5] Alex Kendall and Yarin Gal. What uncertainties do we need in bayesian deep learning for
264 computer vision? In *Advances in Neural Information Processing Systems*, 2017.
- 265 [6] Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances
266 for uncertainty estimation in natural language generation. In *International Conference on*
267 *Learning Representations (ICLR)*, 2023.
- 268 [7] An Yang, Anfeng Yang, Baosong Yang, et al. Qwen3 technical report. *arXiv preprint*
269 arXiv:2505.09388, 2025.
- 270 [8] Zhengyang Yuan, Jiawei Liu, Boyuan Zi, Hao Ning, Kai Zheng, Minghao Chen, et al. Scalequest:
271 Scalable data synthesis for mathematical reasoning. In *International Conference on Learning*
272 *Representations (ICLR)*, 2025.
- 273 [9] AI-MO Team. Numinamath-cot: A large-scale mathematical reasoning dataset with
274 chain-of-thought annotations. 2024. [https://huggingface.co/datasets/AI-MO/](https://huggingface.co/datasets/AI-MO/NuminaMath-CoT)
275 NumaMath-CoT.
- 276 [10] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn
277 Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset.
278 In *Advances in Neural Information Processing Systems*, 2021.
- 279 [11] Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma,
280 Liang Chen, Runxin Xu, et al. Omni-math: A universal olympiad level mathematic benchmark
281 for large language models. *arXiv preprint arXiv:2410.07985*, 2024.
- 282 [12] Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu,
283 Xu Han, Yujie Huang, Yuxiang Zhang, et al. Olympiadbench: A challenging benchmark for
284 promoting agi with olympiad-level bilingual multimodal scientific problems. *arXiv preprint*
285 arXiv:2402.14008, 2024.
- 286 [13] Meituan Longcat Team. Amo-bench: Assessing mathematical olympiad problem solving for
287 large language models. 2025. [https://huggingface.co/datasets/meituan-longcat/](https://huggingface.co/datasets/meituan-longcat/AMO-Bench)
288 AMO-Bench.
- 289 [14] Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, and
290 Dixin Jiang. Wizardlm: Empowering large language models to follow complex instructions. In
291 *International Conference on Learning Representations (ICLR)*, 2024.

- 292 [15] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version
 293 of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- 294 [16] Yoon Kim and Alexander M Rush. Sequence-level knowledge distillation. *arXiv preprint
 295 arXiv:1606.07947*, 2016.
- 296 [17] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, and
 297 Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model.
 298 *Advances in Neural Information Processing Systems*, 36, 2023.

299 **A Implementation Details**

300 **A.1 Distributed Training Architecture**

301 Our implementation uses Ray [?] for distributed actor management and vLLM [?] for efficient
 302 LLM serving across $4 \times$ NVIDIA H100 GPUs. The GPU allocation differs between the two phases.

303 **Phase 1: Offline semantic entropy computation.** This phase runs as a standalone script before
 304 training begins. Two vLLM model instances are loaded on the same GPU. The **teacher vLLM**
 305 **engine** (GPU 0) serves the teacher model (e.g., Qwen3-8B) with `gpu_memory_utilization=0.65`
 306 and `tensor_parallel_size=1`; for each prompt it generates $N=8$ responses in a single batched
 307 call via `engine.generate()`. The **cluster judge vLLM engine** (GPU 0, colocated) serves the
 308 semantic equivalence judge (Qwen2.5-3B-Instruct) with `gpu_memory_utilization=0.20`; it is
 309 only invoked for answer pairs that do not match via exact string comparison, keeping its utilization
 310 low. Both engines share GPU 0 via vLLM’s memory pre-allocation. The remaining GPUs are idle
 311 during this phase. Processing 8,800 prompts with batch size 300 completes in approximately 1–2
 312 hours.

313 **Phase 2: On-policy distillation.** All 4 GPUs are utilized via Ray placement groups. **GPU 0** hosts
 314 the *student actor*: the student model (e.g., Qwen3-1.7B) wrapped in DeepSpeed ZeRO Stage 1
 315 for gradient computation and parameter updates; a frozen reference copy of the initial student is
 316 colocated on the same GPU using fractional allocation (`num_gpus=0.2`). **GPU 1** hosts the *teacher
 317 actor*: the frozen teacher model loaded in inference mode, which receives student-generated rollouts
 318 and returns top- k logit values and token indices at every position. **GPUs 2–3** host two LLMRayActor
 319 *vLLM generation engines*, each initialized from the student’s current weights. They generate $K=4$
 320 on-policy rollouts per prompt in parallel. After each gradient step, the updated student weights are
 321 broadcast to both engines via `update_weight()` over the Gloo backend.

322 The training loop proceeds as follows at each step:

- 323 1. The Ray coordinator dispatches a mini-batch of prompts to the 2 vLLM engines.
- 324 2. Each engine generates $K/2=2$ rollouts per prompt (total $K=4$ per prompt across both
 325 engines).
- 326 3. Rollouts are sent to the teacher actor, which computes top- k logits for all token positions.
- 327 4. The student actor receives rollouts, teacher logits, and precomputed w_{se} weights, then
 328 performs a gradient update.
- 329 5. Updated student weights are synchronized back to both vLLM engines for the next iteration.

330 **A.2 Memory Optimization Techniques**

331 **Top-K Teacher Logits.** Instead of storing full vocabulary logits ($151,936 \times 4$ bytes = 608 KB per
 332 position), we keep only top-512 values and indices:

```
333 topk_vals, topk_ids = logits.topk(k=512, dim=-1)
334 # Store: 512 × 4 bytes (vals) + 512 × 4 bytes (ids) = 4 KB
```

335 This achieves 150x memory reduction with minimal accuracy loss.

336 **Token-by-token Distillation Loss.** We compute KL divergence incrementally to avoid materializing large tensors:

```
338 for t in range(seq_len):
339     teacher_probs_t = F.softmax(teacher_vals[t], dim=-1)
340     student_logprobs_t = F.log_softmax(student_logits[t], dim=-1)
341     student_logprobs_topk = student_logprobs_t.gather(-1, teacher_ids[t])
342     kl_t = -(teacher_probs_t * student_logprobs_topk).sum(-1)
```

343 **A.3 Communication Backend Selection**

344 For non-colocated setups, vLLM workers require collective communication for weight updates.
345 **NCCL** is fast but fails when actors reside in separate Ray placement groups due to
346 **CUDA_VISIBLE_DEVICES** isolation. We therefore use **Gloo**, a CPU-based fallback that works across
347 placement groups, when **tensor_parallel_size=1**.

348 **B Additional Experimental Results**

349 **B.1 Hyperparameter Sensitivity**

350 **TODO:** Add figures/tables for learning rate sweep, temperature sweep for generation, beta sweep for
351 contrastive loss, and number of samples per prompt (K).

352 **B.2 Evaluation Protocol Details**

353 For GSM8K evaluation, we:

- 354 1. Generate greedy responses (temperature=0, n=1)
- 355 2. Extract final numerical answer using regex
- 356 3. Compare with ground-truth using `math_equal()` for numerical equivalence
- 357 4. Report accuracy = correct / total

358 **B.3 Computational Cost**

359 Training on the SURE-Math dataset takes approximately **TODO** wall-clock hours on $4 \times$ H100
360 (**TODO** GPU-hours total). The time breakdown is: on-policy generation **TODO%**, teacher logits
361 computation **TODO%**, and student training **TODO%**.

362 **C Reproducibility**

363 Code will be released at <https://github.com/TODO>. Full training configurations are provided in
364 YAML format.