

---

# Uncertainty-Calibrated On-Policy Distillation for Large Language Models

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

On-policy distillation trains a student on its own outputs with teacher feedback to avoid the train-inference distribution mismatch present in standard distillation. However, a fundamental underexplored challenge is that the teacher itself is imperfect; consequently, naively imitating erroneous teacher outputs propagates mistakes into the student. Prior work such as Self-Distilled Reasoner [1] mitigates this by injecting ground-truth chain-of-thought solutions into the teacher’s system prompt to guarantee the teacher with correct reasoning. But this approach depends on labeled data that is unavailable in many practical settings. We introduce **UOPD** (Uncertainty-calibrated On-Policy Distillation), a framework that addresses teacher errors without ground-truth supervision. Teacher reliability is estimated offline by computing its **uncertainty**, and the distillation objective is reweighted inversely proportional to the teacher’s uncertainty with no additional training cost. We further release **SURE-Math**, a dataset annotated with teacher semantic entropy scores. We evaluate UOPD across mathematical reasoning, summarization, translation, and instruction tuning tasks. UOPD converges faster and achieves state-of-the-art results compared to both GRPO and standard on-policy distillation.

17 **1 Introduction**

18 **2 Preliminary**

19 **2.1 On-Policy Knowledge Distillation**

20 Standard knowledge distillation [2] trains a student model  $p_S(\cdot; \theta_S)$  to minimize the KL divergence  
21 from a frozen teacher  $p_T$  on a fixed dataset  $\mathcal{D}$ :

$$\mathcal{L}_{\text{KD}}(\theta_S) = \mathbb{E}_{x \sim \mathcal{D}} [\text{KL}(p_T(\cdot | x) \| p_S(\cdot | x; \theta_S))]. \quad (1)$$

22 This offline objective suffers from a train-inference distribution mismatch: during training the  
23 student observes token-level distributions conditioned on dataset prefixes, whereas at inference it  
24 must generate from its own previously predicted tokens. On-policy distillation [3] resolves this by  
25 generating training sequences from the student itself:

$$\mathcal{L}_{\text{OPD}}(\theta_S) = \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y \sim p_S(\cdot | x; \theta_S)} \left[ \sum_{t=1}^{|y|} \text{KL}(p_T(\cdot | x, y_{<t}) \| p_S(\cdot | x, y_{<t}; \theta_S)) \right]. \quad (2)$$

26 Since the training distribution now matches the inference distribution, on-policy distillation substantially  
27 reduces compounding errors in chain-of-thought reasoning [4].

28 **2.2 Uncertainty in LLMs**

29 Uncertainty in language model outputs can be decomposed into two fundamentally distinct sources  
30 [5]. **Aleatoric uncertainty** is irreducible and arises from the inherent ambiguity of natural language,  
31 where the same meaning can be expressed in infinitely many surface forms. Consider the prompt  
32 “*What is the capital of the UK?*” Given this question, the responses “*London*”, “*London is the capital*  
33 *of the UK*”, and “*The UK’s capital is London*” are semantically identical, while their token-level  
34 divergence is large. A teacher that consistently produces such paraphrases is not making errors; it  
35 is exhibiting natural lexical variation. Penalising this variation at the token level conflates surface  
36 diversity with genuine unreliability.

37 **Epistemic uncertainty**, by contrast, reflects the teacher’s *lack of knowledge*. It arises when the  
38 teacher generates contradictory answers across samples, indicating that it does not reliably know the  
39 correct response to a given prompt. Standard token-level objectives cannot distinguish these two  
40 sources, because both manifest as high distributional variance in the output space.

41 Semantic entropy [6] resolves this ambiguity by operating over *meaning* rather than surface tokens.  
42 Responses are first clustered into semantic equivalence classes  $\mathcal{C}$ , and entropy is computed over the  
43 resulting distribution

$$H_{\text{sem}}(x) = - \sum_{c \in \mathcal{C}} p(c | x) \log p(c | x), \quad p(c | x) \propto \sum_{y \in c} p_T(y | x). \quad (3)$$

44 Paraphrases of the same answer collapse into a single class, suppressing aleatoric noise.  $H_{\text{sem}}(x)$   
45 is therefore a faithful, label-free measure of epistemic uncertainty; it is high only when the teacher  
46 genuinely disagrees with itself across semantically distinct answers.

47 **3 Method**

48 We propose UOPD, a two-phase framework for uncertainty-calibrated on-policy distillation. In the  
49 first phase, we estimate the teacher’s epistemic uncertainty for each training prompt by computing  
50 semantic entropy over multiple teacher rollouts. This computation is performed entirely offline before  
51 training begins, producing a per-prompt uncertainty weight that is stored alongside the training data.  
52 In the second phase, the student is trained on-policy with token-level distillation from the teacher,  
53 where each sample’s contribution to the loss is calibrated by its precomputed uncertainty weight.

54 The key advantage of this design is the separation of uncertainty estimation from training. Because  
55 semantic entropy is computed once before the training loop, UOPD introduces zero additional  
56 overhead during training while still enabling uncertainty-aware distillation. Figure 1 illustrates the  
57 overall pipeline. We describe each component in detail below.

58 **3.1 Offline Semantic Entropy Estimation**

59 Given a training set of prompts  $\{x_i\}_{i=1}^M$ , we first estimate the teacher’s epistemic uncertainty per  
60 prompt. For each prompt  $x_i$ , we sample  $N$  responses from the teacher  $\{y_1^T, \dots, y_N^T\} \sim p_T(\cdot | x_i)$ ,  
61 recording each response’s sequence log-probability  $\log p_T(y_n^T | x_i)$  and token count  $|y_n^T|$ .

62 **Semantic clustering.** As discussed in Section 2, responses such as “*London*”, “*London is the*  
63 *capital of the UK*”, and “*The UK’s capital is London*” are semantically identical but produce large  
64 token-level divergence. To prevent this aleatoric variation from inflating uncertainty estimates, we  
65 extract the final answer from each teacher response and cluster semantically equivalent answers into  
66 equivalence classes  $\mathcal{C} = \{c_1, \dots, c_J\}$ . Two answers are placed in the same cluster only if one can be  
67 derived from the other, i.e., they express the same meaning in different surface forms (e.g.,  $\frac{1}{2}$  and 0.5,  
68 or  $x = 3$  and 3). Answers that match as exact strings are merged directly. For non-matching pairs, a  
69 Qwen-4B judge [7] determines pairwise semantic equivalence, and all judgments are resolved into  
70 clusters with a Union-Find algorithm.

71 **Intra-cluster vs. inter-cluster entropy.** This clustering decomposes the total entropy of teacher  
72 outputs into two interpretable components. *Intra-cluster entropy* captures the lexical diversity *within*  
73 a semantic equivalence class. A cluster containing “0.5”, “ $\frac{1}{2}$ ”, and “the answer is one half” has high  
74 intra-cluster entropy, reflecting the teacher’s expressive richness rather than genuine confusion. This  
75 variation is aleatoric and should not be penalized. *Inter-cluster entropy*, by contrast, measures the

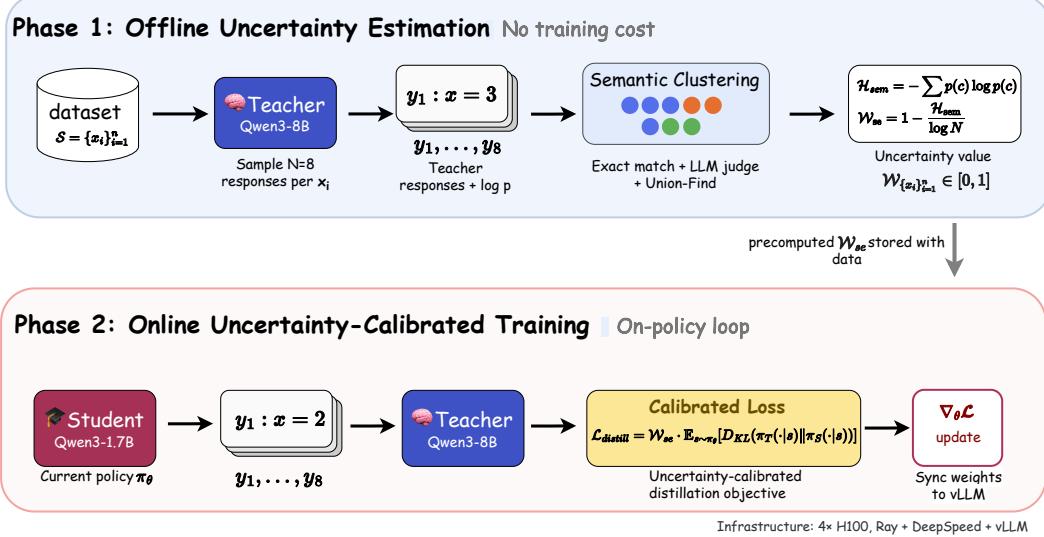


Figure 1: Overview of UOPD. Phase 1 estimates teacher reliability offline by computing semantic entropy over meaning clusters rather than surface-form token distributions. Phase 2 performs on-policy distillation whose per-sample loss is calibrated by the precomputed uncertainty weights. Because the two phases are fully decoupled, uncertainty estimation introduces zero additional overhead during training.

76 spread of probability mass *across* semantically distinct answer classes. When the teacher assigns  
 77 substantial mass to multiple contradictory clusters (e.g., both “3” and “5” for the same prompt),  
 78 inter-cluster entropy is high, signaling epistemic uncertainty. Our semantic entropy  $H_{\text{sem}}$  captures  
 79 precisely the inter-cluster component by collapsing all within-cluster variation before computing  
 80 entropy.

81 **Probability-weighted semantic entropy.** Each response’s log-probability is first length-normalized  
 82 to avoid penalizing longer outputs. The probability mass of semantic class  $c$  is then computed as

$$p(c | x_i) = \frac{\sum_{n \in c} \exp(\log p_T(y_n^T | x_i) / |y_n^T|)}{\sum_{n=1}^N \exp(\log p_T(y_n^T | x_i) / |y_n^T|)}, \quad (4)$$

83 and the semantic entropy is

$$H_{\text{sem}}(x_i) = - \sum_{c \in C} p(c | x_i) \log p(c | x_i). \quad (5)$$

84 We normalize by the maximum possible entropy to obtain a value in  $[0, 1]$  and define the per-prompt  
 85 distillation weight as

$$w_{\text{se}}(x_i) = 1 - \frac{H_{\text{sem}}(x_i)}{\log N}. \quad (6)$$

86 A weight near 1 indicates that the teacher consistently agrees on the same semantic answer (low  
 87 epistemic uncertainty), while a weight near 0 indicates contradictory responses across clusters (high  
 88 epistemic uncertainty). These weights are precomputed and stored with the training data.

### 89 3.2 Uncertainty-Calibrated On-Policy Distillation

90 At each training step, the student generates  $K$  responses per prompt from its current policy  
 91  $\{y^{(1)}, \dots, y^{(K)}\} \sim p_S(\cdot | x; \theta_S)$ . For each student-generated response  $y$ , the frozen teacher pro-  
 92 vides its top- $k$  logit values and corresponding token indices at every position, reducing memory and  
 93 communication from the full vocabulary size  $V$  to  $k \ll V$ .

94 The token-level distillation loss at position  $t$  is the cross-entropy between the teacher’s and student’s  
 95 distributions over the top- $k$  tokens

$$\ell_t = - \sum_{j=1}^k \tilde{p}_T(v_j | x, y_{<t}) \log p_S(v_j | x, y_{<t}; \theta_S), \quad (7)$$

96 where  $\tilde{p}_T$  denotes the teacher’s probability renormalized over the top- $k$  tokens and  $\{v_1, \dots, v_k\}$  are  
 97 the teacher’s top- $k$  token indices.

98 The standard on-policy distillation objective averages  $\ell_t$  uniformly over all response tokens, treating  
 99 every prompt equally regardless of teacher reliability. UOPD instead calibrates each sample’s  
 100 contribution by its precomputed semantic entropy weight  $w_{se}(x)$ . For a prompt  $x$  and student-  
 101 generated response  $y$ , the uncertainty-calibrated loss is

$$\mathcal{L}(x, y; \theta_S) = \frac{w_{se}(x)}{\sum_t \mathbb{1}[t \in \text{resp}]} \sum_{t \in \text{resp}} \ell_t, \quad (8)$$

102 where the sum runs over response token positions (excluding the prompt). When the teacher is  
 103 confident about a prompt ( $w_{se} \approx 1$ ), the student receives the full distillation signal. When the teacher  
 104 is uncertain ( $w_{se} \approx 0$ ), the loss is suppressed, preventing the propagation of erroneous teacher  
 105 guidance.

### 106 3.3 Overall Training Objective

107 The full training objective averages over all prompts and their  $K$  on-policy rollouts

$$\mathcal{L}_{\text{UOPD}}(\theta_S) = \frac{1}{|\mathcal{D}|} \sum_{x \in \mathcal{D}} \frac{1}{K} \sum_{j=1}^K \mathcal{L}(x, y^{(j)}; \theta_S). \quad (9)$$

108 At each training step, the updated student weights are synchronized back to the vLLM generation  
 109 engines, ensuring that the next round of rollouts reflects the latest policy. This on-policy loop  
 110 continues until convergence.

111 Algorithm 1 summarizes the complete UOPD pipeline. Phase 1 (lines 2–7) iterates over all training  
 112 prompts, sampling  $N$  teacher responses per prompt, clustering them into semantic equivalence classes,  
 113 and storing the resulting uncertainty weight  $w_{se}$ . Phase 2 (lines 9–14) performs the on-policy training  
 114 loop: at each step, the student generates  $K$  rollouts, queries the teacher for top- $k$  logits at every token  
 115 position, and updates its parameters with the uncertainty-weighted distillation loss, after which the  
 116 refreshed weights are synchronized to the generation engines for the next iteration.

## 117 4 Experiments

### 118 4.1 SURE-Math Dataset

119 We introduce **SURE-Math** (Semantic-Uncertainty REasoning Math), a mathematical reasoning  
 120 dataset in which every problem is annotated with the teacher’s precomputed semantic entropy score.

121 We first collect 1,000 mathematics problems spanning middle-school to competition level, each paired  
 122 with a verified ground-truth answer. Of these, 200 are reserved as a held-out test set (Section 4.3); the  
 123 remaining 800 are added to the training pool.

124 We further aggregate seed problems from public sources across four difficulty tiers. *Easy*: ScaleQuest-  
 125 Math [8] (15K). *Medium*: NuminMath-CoT [9] (10K) and MATH [10] (12.5K). *Hard*: Omni-MATH  
 126 [11] (4.4K) and OlympiadBench [12] (5K). *Competition*: AIME 2024–2025, HMMT Feb & Nov  
 127 2025, and AMO-Bench [13].

128 Starting from these seeds, we synthesize new problems with Qwen2.5-72B-Instruct using an Evol-  
 129 Instruct [14] style pipeline. Six evolution strategies are applied, each realized by a dedicated  
 130 system prompt: *harder* (increase reasoning steps or add constraints), *rewrite* (change context and  
 131 wording while preserving the underlying skill), *algebraize* (replace concrete values with variables),  
 132 *apply* (embed the concept in a real-world scenario), *compose* (combine with a different branch of

---

**Algorithm 1** Uncertainty-Calibrated On-Policy Distillation (UOPD)

---

**Require:** Training prompts  $\mathcal{D} = \{x_i\}_{i=1}^M$ , teacher  $p_T$ , student  $p_S(\cdot; \theta_S)$ , SE samples  $N$ , rollouts  $K$ , top- $k$

1: — **Phase 1: Offline Semantic Entropy (Sec. 3.1) —**

2: **for** each prompt  $x_i \in \mathcal{D}$  **do**

3:   Sample  $N$  responses  $\{y_1^T, \dots, y_N^T\} \sim p_T(\cdot | x_i)$

4:   Extract final answers; cluster into semantic classes  $\mathcal{C}$

5:   Compute semantic entropy  $H_{\text{sem}}(x_i)$  (Eq. 5)

6:   Store weight  $w_{\text{se}}(x_i) = 1 - H_{\text{sem}}(x_i) / \log N$  (Eq. 6)

7: **end for**

8: — **Phase 2: On-Policy Distillation (Sec. 3.2) —**

9: **while** not converged **do**

10:   Sample mini-batch  $\mathcal{B} \subset \mathcal{D}$

11:   **for** each  $x \in \mathcal{B}$  **do**

12:     Generate  $K$  rollouts  $\{y^{(1)}, \dots, y^{(K)}\} \sim p_S(\cdot | x; \theta_S)$

13:     **for** each rollout  $y^{(j)}$  **do**

14:       Query teacher for top- $k$  logits at each token position of  $y^{(j)}$

15:       Compute token-level CE:  $\ell_t = -\sum_v \tilde{p}_T(v) \log p_S(v; \theta_S)$

16:     **end for**

17:     Compute weighted loss  $\mathcal{L}(x, y^{(j)}; \theta_S)$  (Eq. 8)

18:   **end for**

19:   Update  $\theta_S$  with  $\nabla_{\theta_S} \frac{1}{|\mathcal{B}|} \sum_{x \in \mathcal{B}} \frac{1}{K} \sum_{j=1}^K \mathcal{L}(x, y^{(j)}; \theta_S)$

20:   Sync  $\theta_S$  to vLLM generation engines

21: **end while**

---

133 mathematics), and *competition* (transform into AMC/AIME/Olympiad style). Each seed undergoes  
134 1–5 rounds of iterative evolution depending on its difficulty tier: easy seeds receive 1 round, medium  
135 seeds 2 rounds, hard seeds 3 rounds, and competition seeds 5 rounds. In follow-up rounds, a different  
136 strategy (drawn from *harder*, *competition*, and *compose*) is applied to the output of the previous  
137 round, progressively increasing difficulty. Strategy selection within each tier is weighted to favor  
138 *harder* and *competition* for higher-tier seeds.

139 For each of the 8,000 problems, we sample  $N=8$  responses from Qwen3-8B [7] at temperature 0.7,  
140 extract the \boxed{} answer from each response, and compute the probability-weighted semantic  
141 entropy as described in Section 3.1. The resulting per-problem semantic entropy score and the  
142 corresponding distillation weight  $w_{\text{se}}$  are stored alongside each problem.

143 We reserve 200 of the 1,000 manually collected problems (which carry ground-truth labels) as a  
144 held-out test set. The remaining 800 curated problems and all 8,000 evolved problems form the  
145 training set.

146 **4.2 Experimental Setup**

147 **Models.** We evaluate three student–teacher pairs of increasing scale: Qwen3-1.7B/Qwen3-8B,  
148 Qwen3-4B/Qwen3-8B, and Qwen3-8B/Qwen3-30B [7], all instruction-tuned variants. Semantic  
149 clustering for the offline SE computation uses a Qwen2.5-3B-Instruct judge for pairwise semantic  
150 equivalence.

151 **Training details.** We train with the on-policy distillation objective described in Eq. 8. At each step,  
152 the student generates  $K=4$  rollouts per prompt. The teacher provides its top-512 logit values at every  
153 token position. We use a learning rate of  $3 \times 10^{-6}$  with a cosine schedule and 5% warmup, AdamW  
154 with  $(\beta_1, \beta_2) = (0.9, 0.95)$ , gradient clipping at 1.0, and a global batch size of 128. Training runs for  
155 one epoch over 50,000 samples (with replacement from the 8,800 training problems). All models use  
156 bfloat16 precision and FlashAttention-2.

157 For each student–teacher pair, we compare UOPD against the following methods. (1) **Base**: the  
158 student model without any distillation (lower bound). (2) **Teacher**: the teacher model (upper  
159 bound). (3) **Standard OPD**: on-policy distillation with uniform weighting ( $w_{\text{se}} = 1$  for all prompts),  
160 following [3]. (4) **GRPO**: Group Relative Policy Optimization, which trains the student using reward

Table 1: Main results on mathematical reasoning benchmarks. We report pass@1 (greedy) and avg@16 accuracy (%). Best student-sized results are **bolded**.

Method	MATH-500		AIME 2025		AMO-Bench		SURE-Math	
	pass@1	avg@16	pass@1	avg@16	pass@1	avg@16	pass@1	avg@16
<b>Student: Qwen3-1.7B Teacher: Qwen3-8B</b>								
Base (1.7B)	–	–	–	–	–	–	–	–
Standard OPD	–	–	–	–	–	–	–	–
GRPO	–	–	–	–	–	–	–	–
Self-Distilled Reasoner	–	–	–	–	–	–	–	–
<b>UOPD (Ours)</b>	–	–	–	–	–	–	–	–
<b>Student: Qwen3-4B Teacher: Qwen3-8B</b>								
Base (4B)	–	–	–	–	–	–	–	–
Standard OPD	–	–	–	–	–	–	–	–
GRPO	–	–	–	–	–	–	–	–
Self-Distilled Reasoner	–	–	–	–	–	–	–	–
<b>UOPD (Ours)</b>	–	–	–	–	–	–	–	–
<b>Student: Qwen3-8B Teacher: Qwen3-30B</b>								
Base (8B)	–	–	–	–	–	–	–	–
Standard OPD	–	–	–	–	–	–	–	–
GRPO	–	–	–	–	–	–	–	–
Self-Distilled Reasoner	–	–	–	–	–	–	–	–
<b>UOPD (Ours)</b>	–	–	–	–	–	–	–	–

161 signals from correct and incorrect rollouts. (5) **Self-Distilled Reasoner**: on-policy distillation with  
 162 ground-truth chain-of-thought injected into the teacher’s system prompt [1].

### 163 4.3 Evaluation

164 **Benchmarks.** We evaluate all methods on three public mathematical reasoning benchmarks of  
 165 increasing difficulty, plus our held-out test set: **MATH-500** [10], 500 competition-level problems  
 166 spanning seven subjects; **AIME 2025**, 30 problems from the American Invitational Mathematics  
 167 Examination (Parts I and II); **AMO-Bench** [13], 50 IMO-level competition problems; and our  
 168 held-out **SURE-Math test set** (200 curated problems with ground-truth labels).

169 **Metrics.** We report pass@1 accuracy with greedy decoding (temperature 0) and avg@16 accuracy  
 170 with 16 samples at temperature 1.2 and top- $p=0.95$ .

### 171 4.4 Main Results

172 Table 1 presents the main comparison across all benchmarks.

### 173 4.5 Ablation Studies

174 **Effect of uncertainty weighting.** We compare three weighting strategies: (1) uniform weighting  
 175 ( $w_{se} = 1$ ), which reduces to standard on-policy distillation; (2) binary filtering, which discards all  
 176 prompts with  $H_{sem} > \tau$  for a threshold  $\tau$ ; and (3) soft weighting ( $w_{se} = 1 - H_{sem}/\log N$ ), which is  
 177 our default. Table 2 reports the results.

Table 2: Ablation on uncertainty weighting strategies.

Weighting	MATH-500	AIME 2025	AMO-Bench	SURE-Math
Uniform ( $w = 1$ )	–	–	–	–
Binary filter ( $\tau = 0.3$ )	–	–	–	–
Binary filter ( $\tau = 0.5$ )	–	–	–	–
Soft weighting (UOPD)	–	–	–	–

178 **Number of SE samples**  $N$ . The number of teacher rollouts  $N$  used for semantic entropy estimation  
179 controls the resolution of the uncertainty estimate. We vary  $N \in \{2, 4, 8, 16\}$  and measure  
180 downstream distillation performance.

181 **Top- $k$  logit truncation.** Transmitting the full vocabulary ( $|V| = 151,936$ ) is memory-intensive.  
182 We compare  $k \in \{128, 512, 2048\}$  and full-vocabulary distillation, measuring both accuracy and  
183 peak GPU memory.

184 **Number of student rollouts**  $K$ . We ablate  $K \in \{1, 2, 4, 8\}$  on-policy rollouts per prompt to  
185 understand the trade-off between training diversity and computational cost.

## 186 4.6 Analysis

### 187 Convergence speed.

188 **Semantic entropy distribution.** Figure ?? shows the distribution of semantic entropy scores across  
189 SURE-Math. The majority of problems have low SE, indicating that the teacher is confident on most  
190 prompts. The long tail of high-SE problems represents cases where the teacher genuinely disagrees  
191 with itself, and these are precisely the samples that UOPD downweights.

192 **Qualitative examples.** Table ?? shows representative examples where UOPD’s uncertainty weight-  
193 ing helps. For low-SE prompts, the teacher provides consistent guidance and the student learns  
194 effectively. For high-SE prompts, the teacher produces contradictory answers; UOPD suppresses  
195 these samples, preventing the student from learning incorrect reasoning patterns.

## 196 5 Related Work

### 197 5.1 Knowledge Distillation for Language Models

198 Traditional knowledge distillation [2] trains student models to match teacher output distributions.  
199 For autoregressive language models, supervised KD [15] and sequence-level KD [16] are widely  
200 used. However, these off-policy approaches suffer from distribution mismatch between training  
201 (teacher-generated or ground-truth sequences) and inference (student-generated sequences).

202 **On-policy distillation** [3] addresses this by training students on their own generated outputs, using  
203 teacher logits as labels. GKD (Generalized Knowledge Distillation) demonstrates strong improve-  
204 ments on summarization and translation tasks. Our work extends on-policy distillation to reasoning  
205 tasks by incorporating verification signals and contrastive objectives.

### 206 5.2 Learning from Verification Feedback

207 Recent work on mathematical reasoning leverages verification to improve model training. GRPO [? ]  
208 and similar RL-based approaches optimize for verified correctness using policy gradient methods.  
209 V-STaR [? ] iteratively generates verified solutions for self-improvement.

210 However, these methods rely solely on sparse binary rewards (correct/incorrect) and do not leverage  
211 dense token-level teacher guidance. TCD bridges this gap by combining verification with token-level  
212 distillation.

### 213 5.3 Contrastive Learning for Language Models

214 DPO [17] introduced preference-based contrastive learning for alignment, training models to prefer  
215 chosen responses over rejected ones. SimPO [? ] and other variants explore different contrastive  
216 formulations.

217 Our work adapts contrastive learning to the distillation setting: we use verified correct responses  
218 (from reference data or teacher generation) as “chosen” and student-generated incorrect responses as  
219 “rejected”, enabling the student to learn from its mistakes with teacher guidance.

220 **6 Conclusion**

221 We presented Token-level Contrastive Distillation (TCD), a framework that combines on-policy  
222 generation, verification-based learning, and dense teacher guidance for distilling reasoning capabilities  
223 into smaller language models. By unifying distillation on correct traces with contrastive learning  
224 on errors, TCD effectively leverages both sparse verification signals and rich token-level teacher  
225 feedback.

226 Our experiments on GSM8K demonstrate that TCD enables a 1.7B student model to achieve strong  
227 mathematical reasoning performance when distilled from an 8B teacher. The framework is efficient  
228 and scalable, using vLLM for on-policy generation and memory-optimized techniques for handling  
229 large vocabularies.

230 **Limitations and Future Work.** **Scalability:** Current experiments use 4 GPUs; scaling to multi-  
231 node setups requires careful optimization of communication backends (NCCL vs. Gloo). **Generaliza-**  
232 **tion:** We focus on mathematical reasoning; extending to other reasoning tasks (code, commonsense)  
233 is an important direction. **Teacher quality:** Our approach assumes access to a capable teacher;  
234 exploring self-improvement scenarios is promising.

235 TCD opens avenues for efficient reasoning model deployment by making strong reasoning capabilities  
236 accessible in compact models suitable for resource-constrained environments.

237 **References**

- 238 [1] Siyan Zhao, Zhihui Xie, Mengchen Liu, Jing Huang, Guan Pang, Feiyu Chen, and Aditya  
239 Grover. Self-distilled reasoner: On-policy self-distillation for large language models. *arXiv*  
240 preprint arXiv:2601.18734, 2026.
- 241 [2] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network.  
242 *arXiv preprint arXiv:1503.02531*, 2015.
- 243 [3] Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu  
244 Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-  
245 generated mistakes. In *International Conference on Learning Representations (ICLR)*, 2024.
- 246 [4] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, and  
247 Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. *Advances*  
248 in *Neural Information Processing Systems*, 35:24824–24837, 2022.
- 249 [5] Alex Kendall and Yarin Gal. What uncertainties do we need in bayesian deep learning for  
250 computer vision? In *Advances in Neural Information Processing Systems*, 2017.
- 251 [6] Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances  
252 for uncertainty estimation in natural language generation. In *International Conference on*  
253 *Learning Representations (ICLR)*, 2023.
- 254 [7] An Yang, Anfeng Yang, Baosong Yang, et al. Qwen3 technical report. *arXiv preprint*  
255 arXiv:2505.09388, 2025.
- 256 [8] Zhengyang Yuan, Jiawei Liu, Boyuan Zi, Hao Ning, Kai Zheng, Minghao Chen, et al. Scalequest:  
257 Scalable data synthesis for mathematical reasoning. In *International Conference on Learning*  
258 *Representations (ICLR)*, 2025.
- 259 [9] AI-MO Team. Numinamath-cot: A large-scale mathematical reasoning dataset with  
260 chain-of-thought annotations. 2024. [https://huggingface.co/datasets/AI-MO/](https://huggingface.co/datasets/AI-MO/NuminaMath-CoT)  
261 Numinamath-CoT.
- 262 [10] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn  
263 Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset.  
264 In *Advances in Neural Information Processing Systems*, 2021.
- 265 [11] Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma,  
266 Liang Chen, Runxin Xu, et al. Omni-math: A universal olympiad level mathematic benchmark  
267 for large language models. *arXiv preprint arXiv:2410.07985*, 2024.

- 268 [12] Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu,  
 269 Xu Han, Yujie Huang, Yuxiang Zhang, et al. Olympiadbench: A challenging benchmark for  
 270 promoting agi with olympiad-level bilingual multimodal scientific problems. *arXiv preprint*  
 271 *arXiv:2402.14008*, 2024.
- 272 [13] Meituan Longcat Team. Amo-bench: Assessing mathematical olympiad problem solving for  
 273 large language models. 2025. <https://huggingface.co/datasets/meituan-longcat/AMO-Bench>.
- 275 [14] Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, and  
 276 Dixin Jiang. Wizardlm: Empowering large language models to follow complex instructions. In  
 277 *International Conference on Learning Representations (ICLR)*, 2024.
- 278 [15] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version  
 279 of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- 280 [16] Yoon Kim and Alexander M Rush. Sequence-level knowledge distillation. *arXiv preprint*  
 281 *arXiv:1606.07947*, 2016.
- 282 [17] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, and  
 283 Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model.  
 284 *Advances in Neural Information Processing Systems*, 36, 2023.

## 285 A Implementation Details

### 286 A.1 Distributed Training Architecture

287 Our implementation uses Ray [?] for distributed actor management and vLLM [?] for efficient  
 288 LLM serving across 4× NVIDIA H100 GPUs. The GPU allocation differs between the two phases.

289 **Phase 1: Offline semantic entropy computation.** This phase runs as a standalone script before  
 290 training begins. Two vLLM model instances are loaded on the same GPU. The **teacher vLLM**  
 291 **engine** (GPU 0) serves the teacher model (e.g., Qwen3-8B) with `gpu_memory_utilization=0.65`  
 292 and `tensor_parallel_size=1`; for each prompt it generates  $N=8$  responses in a single batched  
 293 call via `engine.generate()`. The **cluster judge vLLM engine** (GPU 0, colocated) serves the  
 294 semantic equivalence judge (Qwen2.5-3B-Instruct) with `gpu_memory_utilization=0.20`; it is  
 295 only invoked for answer pairs that do not match via exact string comparison, keeping its utilization  
 296 low. Both engines share GPU 0 via vLLM’s memory pre-allocation. The remaining GPUs are idle  
 297 during this phase. Processing 8,800 prompts with batch size 300 completes in approximately 1–2  
 298 hours.

299 **Phase 2: On-policy distillation.** All 4 GPUs are utilized via Ray placement groups. **GPU 0**  
 300 hosts the *student actor*: the student model (e.g., Qwen3-1.7B) wrapped in DeepSpeed ZeRO Stage 1  
 301 for gradient computation and parameter updates; a frozen reference copy of the initial student is  
 302 colocated on the same GPU using fractional allocation (`num_gpus=0.2`). **GPU 1** hosts the *teacher*  
 303 *actor*: the frozen teacher model loaded in inference mode, which receives student-generated rollouts  
 304 and returns top- $k$  logit values and token indices at every position. **GPUs 2–3** host two LLMRayActor  
 305 *vLLM generation engines*, each initialized from the student’s current weights. They generate  $K=4$   
 306 on-policy rollouts per prompt in parallel. After each gradient step, the updated student weights are  
 307 broadcast to both engines via `update_weight()` over the Gloo backend.

308 The training loop proceeds as follows at each step:

- 309 1. The Ray coordinator dispatches a mini-batch of prompts to the 2 vLLM engines.
- 310 2. Each engine generates  $K/2=2$  rollouts per prompt (total  $K=4$  per prompt across both  
 311 engines).
- 312 3. Rollouts are sent to the teacher actor, which computes top- $k$  logits for all token positions.
- 313 4. The student actor receives rollouts, teacher logits, and precomputed  $w_{se}$  weights, then  
 314 performs a gradient update.
- 315 5. Updated student weights are synchronized back to both vLLM engines for the next iteration.

316 **A.2 Memory Optimization Techniques**

317 **Top-K Teacher Logits.** Instead of storing full vocabulary logits ( $151,936 \times 4$  bytes = 608 KB per  
318 position), we keep only top-512 values and indices:

```
319 topk_vals, topk_ids = logits.topk(k=512, dim=-1)
320 # Store: 512 × 4 bytes (vals) + 512 × 4 bytes (ids) = 4 KB
```

321 This achieves 150× memory reduction with minimal accuracy loss.

322 **Token-by-token Distillation Loss.** We compute KL divergence incrementally to avoid materializing  
323 large tensors:

```
324 for t in range(seq_len):
325     teacher_probs_t = F.softmax(teacher_vals[t], dim=-1)
326     student_logprobs_t = F.log_softmax(student_logits[t], dim=-1)
327     student_logprobs_topk = student_logprobs_t.gather(-1, teacher_ids[t])
328     kl_t = -(teacher_probs_t * student_logprobs_topk).sum(-1)
```

329 **A.3 Communication Backend Selection**

330 For non-colocated setups, vLLM workers require collective communication for weight up-  
331 dates. NCCL is fast but fails when actors reside in separate Ray placement groups due to  
332 CUDA\_VISIBLE\_DEVICES isolation. We therefore use **Gloo**, a CPU-based fallback that works across  
333 placement groups, when tensor\_parallel\_size=1.

334 **B Additional Experimental Results**

335 **B.1 Hyperparameter Sensitivity**

336 **TODO:** Add figures/tables for learning rate sweep, temperature sweep for generation, beta sweep for  
337 contrastive loss, and number of samples per prompt ( $K$ ).

338 **B.2 Evaluation Protocol Details**

339 For GSM8K evaluation, we:

- 340 1. Generate greedy responses (temperature=0, n=1)
- 341 2. Extract final numerical answer using regex
- 342 3. Compare with ground-truth using math\_equal() for numerical equivalence
- 343 4. Report accuracy = correct / total

344 **B.3 Computational Cost**

345 Training on the SURE-Math dataset takes approximately **TODO** wall-clock hours on  $4 \times$  H100  
346 (**TODO** GPU-hours total). The time breakdown is: on-policy generation **TODO**%, teacher logits  
347 computation **TODO**%, and student training **TODO**%.

348 **C Reproducibility**

349 Code will be released at <https://github.com/TODO>. Full training configurations are provided in  
350 YAML format.